

# Pioneer *AWS* Linux Training

# Housekeeping

- THERE ARE NO BAD QUESTIONS
- Intros
  - Name
  - Experience
- Remember there is almost always more than one way to do something

# Prerequisites (if you want to follow along)

- An installation of Mint/Ubuntu/Debian (macOS will work as well)
- Access to a bash shell

# Some basics

- What is Linux?
  - Linux is just the kernel, the interface to the hardware
  - A distribution (e.g. Ubuntu / Mint) is all the stuff that we use to do work packaged up conveniently (more or less)
- The file system starts at / (and it is all downhill from there)
- /var is (generally) where the system writes logs
- /etc is (generally) where config files are kept
- /usr is (generally) where libraries and applications are stored
- /proc is a funny imaginary directory with a bunch of cool stuff like a directory for every process

# Navigating basics

- `ls [directory or file]`
  - List files in a directory
- `ls -lah [directory or file]`
  - Common usage details; all files; human readable
- `ls -ltr [directory or file]`
  - Common usage details; sort by time; reverse the sort order
- `cd [directory]`
  - Change directory
- `pwd`
  - Prints the current working directory
- The shortcut `.` references the current directory
- The shortcut `..` references the parent directory
- The shortcut `~` references your home directory

# Navigating knowledge check

- What command(s) would you use to
  - move to the directory one level up and list the contents?
  - list the contents of your home directory?

# Navigating knowledge check (my way)

- What command(s) would you use to
  - move to the directory one level up and list the contents?  
`cd ..; ls`
  - list the contents of your home directory?  
`ls -ltr ~`

# File permissions

- Every file has three groups of three permissions
  - What everybody can do to the file (or directory)
  - What the file (or directory) group can do
  - What the file (or directory) owner can do
- The permissions in the groups have letters and a “score”
  - read=r (4 points)
  - write=w (2 points)
  - execute=x (1 point)
- Other attributes of a file or directory
  - Type: directory=d; regular=-; symbolic link=l
  - Owner: who owns the file or directory
  - Group: what group of users are associated with the file



# File permissions (continued)

- The permissions can be changed with chmod
  - For each of the three groups just add up the score of what you want it to look like
  - Example: you want a file to be readable and executable by only you,  $4 + 1 = 5$ 
    - `chmod 500 something.sh`
  - Example: you want a file read/write for yourself and readable for everybody else;  $4 + 2 = 6$ ,  $4 = 4$ ,  $4 = 4$ 
    - `chmod 644 another.txt`
  - Any usable directory needs the executable 1
- Other attributes of a file or directory
  - Type: directory=d; regular=-; symbolic link=l
  - Owner: who owns the file or directory
  - Group: what group of users are associated with the file

# Permissions knowledge check

- How would you
  - make a file read only by the current user?
  - make a file readable by all users and writeable by only the owner?
  - make a file readable, writable, and executable by everyone?

# Permissions knowledge check (my way)

- How would you

- make a file read only by the current user?

```
chmod 400 somefile.txt
```

- make a file readable by all users and writeable by only the owner?

```
chmod 644 another.txt
```

- make a file readable, writable, and executable by everyone?

```
chmod 777 harmless.sh
```

# File system manipulation

- **cp** *<source>* *<destination>*
  - Copies a file from one location to another. The source file remains unaltered. More than one source can be listed. Works with wild cards (\*)
  - Use `-r` for moving directories
- **mv** *<source>* *<destination>*
  - Similar to copy but moves the source file. Think of it as a Cut operation in windows. More than one source can be listed. Works with wild cards (\*)
- **rm** *<some-file>*
  - Removes the file listed as the target. More than one target can be listed. Works with wild cards (\*)
  - `-r` for directories. `-f` for read only files (BE CAREFUL!)
- **mkdir** *<new-directory>*
  - Makes a new directory at the target path
  - `mkdir -p <directory-tree>` makes a stack of directories
    - e.g. `mkdir -p a/b/c`

# Viewing file contents

- `cat <some-file>`
  - Outputs the contents of a file to the current shell
- `less <some-file>`
  - Views file contents in current shell and supports paging. More is old but still used heavily, less is an updated version of more
- `tail <some-file>`
  - Used to view the last 10 lines of a file.
- `tail -n 23 <some-file>`
  - See the last 23 lines of the file
- `head <some-file>`
  - See the first 10 lines of a file
- `head -n 23 <some-file>`
  - See the first 23 lines of a file

# File knowledge check

- How would you:
  - Copy a file from directory `a` to directory `b` and then delete the file in directory `a`?
  - View the last 50 lines of a file?
  - Move all files from directory `a` that have a `.txt` extension to directory `b`?
  - Print all lines of a file to the console?

# File knowledge check (my way)

- How would you:

- Copy a file from directory a to directory b and then delete the file in directory a?

```
cp a/somefile.txt b/.
```

- View the last 50 lines of a file?

```
tail -50 b/somefile.txt
```

- Move all files from directory a that have a .txt extension to directory b?

```
mv a/*.txt b/.
```

- Print all lines of a file to the console?

```
cat b/somefile.txt
```

# Files and input

- `echo "some text"`
  - Prints the provided text back to the screen
  - Can also print variables when executed in a script or from environment
- `touch <some-file>`
  - Creates an empty file at the target path or updates the modification time of an existing file
- `something.sh > output.txt 2> error.txt`
  - The `>` operator writes the output of the command to the file listed overwriting the current file contents
  - The `"2"` indicates `stderr` for `something.sh`
- `something.sh >> output.txt`
  - The `>>` operator appends the output of the command to the file listed



# Files and input knowledge check

- How would you:
  - Write the text “this is my sample text” to a new file?
  - Write the text “this is my appended sample text” so that it is appended to an existing file?
  - View the contents of the appended text file?

# Files and input knowledge check (my way)

- How would you:

- Write the text “this is my sample text” to a new file?

```
echo "this is my sample text" > newfile.txt
```

- Write the text “this is my appended sample text” so that it is appended to an existing file?

```
echo "this my appended sample text" >> newfile.txt
```

- View the contents of the appended text file?

```
cat newfile.txt
```

# Other utilities

- *command | command*
  - The | (pipe) operator sends the output of the first command as input to the second command
- **sort**
  - Sorts lines of input file or input sent via pipe
- **wc**
  - Prints number of lines, number of words, and characters
- **env**
  - Prints all environment variables for the current shell
- **ps**
  - Displays a list of all running processes for current user. Add aux to see all processes
- **wget** a “simple” file downloader
  - `wget http://www.google.com`
- **curl** used to make HTTP request from the shell
  - `curl -v http://www.google.com > google.html`

# Other utilities continued

- **top**
  - See running processes and server resource utilization
- **tar**
  - Used to extract or create tar files.
  - Extract .gz file: `tar xzvf filepath`
  - Extract .bz2 file: `tar xjvf filepath`
  - Create .gz file: `tar cvzf compressedfilepath pathfiles`
- **grep**
  - Used to search for text in files or input text
- **Some editor and there are many choices. Learning `vi` or `emacs` is way beyond the scope of this little talk**
  - `nano` is pretty simple text editor
  - `vi` is VERY powerful but hard to learn
  - `emacs` is also VERY powerful and differently hard to learn (note: I've never used `emacs`)

# Utilities knowledge check

- How would you:
  - Get the count of the number of processes running on the system?
  - Get the HTTP response from google.com?
  - List the current environment variables, sort them in alphabetical order, and then write to a local file?
  - Download a file from a remote server and then find a specific string in the file contents? How would you do this in one line?

# Utilities knowledge check (my way)

- How would you:

- Get the count of the number of processes running on the system?

```
ps auxwww | wc -l
```

- Get the HTTP response from google.com?

```
curl -v http://www.google.com > /dev/null
```

- List the current environment variables, sort them in alphabetical order, and then write to a local file?

```
env | sort > sorted_env.txt
```

- Download a file from a remote server and then find a specific string in the file contents? How would you do this in one line?

```
wget http://www.ipchicken.com | grep IP
```

# Package Managers

- Used to install software from remote repositories and make software installs easier
- Ubuntu, Debian, Mint
  - apt-get
- Language specific
  - Python - pip

# Environment file

- In your home directory `.bashrc` contains all the environment variables for your shell that you can set at the start
  - Run on new shell/terminal
- `/etc/profile` has a bunch of stuff is set for all users



# Profile and rc files – example settings

- `alias simple_cmd='complex_cmd'`
  - Used to create command shortcuts
- `export name='value'`
  - Used to set environment variables accessible by child processes
  - Environment variables can be referenced by prepending a \$ character
  - Example - appending and prepending to existing variables
    - `export PATH=/some/addtion:$PATH`
    - `export PATH=$PATH:/some/addition`

# .bashrc knowledge check

- How would you:
  - Create an alias to monitor the file `/var/log/system.log` as a new lines are written to the log file?
  - Append the path `/usr/bin` to the environment variable named `my_env`?

# .bashrc knowledge check (my way)

- How would you:

- Create an alias to monitor the file `/var/log/system.log` as a new lines are written to the log file?

```
echo "alias watchlog='tail -f /var/log/system.log'"  
>> ~/.bashrc
```

- Append the path `/usr/bin` to the environment variable named `my_env`?

```
export my_env=${my_env}:/usr/bin
```

# SSH, SSH Config, SCP

- `ssh remote_user@remote_server`
  - Usually a SSH key is required. The `-i` option facilitates this
  - Private SSH keys typically stored under `~/.ssh/`
  - Keys must only be accessible to user (`chmod 600`)
  - `-L` used for local forwarding
  - `-D` used for dynamic forwarding
  - Make it simple with SSH config...
- `~/.ssh/config`
  - Preconfigure connections into simple alias

# SCP

- *scp source\_server:source\_path remote\_server:remote\_path*
  - Used to copy files between systems
  - if source\_server or remote\_server are not provided, localhost is used by default
  - Again, simple with SSH config

# Job Control

- Processes you start can run (basically) in 3 states
  - foreground: in control of the shell
  - background: the process is running in the background and you can use the shell
    - e.g. `long_runner.sh \>output.txt 2> error.txt &`
  - backgrounded and detached: with the `nohup` command you can put a really long process into the background and exit the shell not killing the process. Output will go to `nohup.out`
    - e.g. `nohup really_long_runner.sh &`
- Signals can be sent with `ctl-c` (kill) and `ctl-z` (stop until `fg`)

# The End

- There is a lot to linux and this just barely scratched the surface
- Let me (or others) help as needed
- Google knows way more about everything than I do
- Questions?